

Natural Language Processing

Tel Aviv University

Assignment 2: Language ModelsDue Date: *Tuesday, Feb 27th EOD, 2024*

Lecturer: Maor Ivgi

1 Word-Level Neural Bi-gram Language Model

- (a) Derive the gradient with respect to the input of a softmax function when cross entropy loss is used for evaluation, i.e., find the gradients with respect to the softmax input vector θ , when the prediction is made by $\hat{y} = \text{softmax}(\theta)$. Cross entropy and softmax are defined as:

$$\text{CE}(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_i y_i \cdot \log(\hat{y}_i)$$

$$\text{softmax}(\theta)_i = \frac{\exp(\theta_i)}{\sum_j \exp(\theta_j)}$$

The gold vector \mathbf{y} is a one-hot vector, and the predicted vector $\hat{\mathbf{y}}$ is a probability distribution over the output space.

- (b) Derive the gradients with respect to the input \mathbf{x} in a one-hidden-layer neural network (i.e., find $\frac{\partial J}{\partial \mathbf{x}}$, where J is the cross entropy loss $\text{CE}(\mathbf{y}, \hat{\mathbf{y}})$). The neural network employs a sigmoid activation function for the hidden layer, and a softmax for the output layer. Assume a one-hot label vector \mathbf{y} is used. The network is defined as:

$$\mathbf{h} = \sigma(\mathbf{x}\mathbf{W}_1 + \mathbf{b}_1),$$

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{h}\mathbf{W}_2 + \mathbf{b}_2).$$

The dimensions of the vectors and matrices are $\mathbf{x} \in \mathbb{R}^{1 \times D_x}$, $\mathbf{h} \in \mathbb{R}^{1 \times D_h}$, $\hat{\mathbf{y}} \in \mathbb{R}^{1 \times D_y}$, $\mathbf{y} \in \mathbb{R}^{1 \times D_y}$. The dimensions of the parameters are $\mathbf{W}_1 \in \mathbb{R}^{D_x \times D_h}$, $\mathbf{W}_2 \in \mathbb{R}^{D_h \times D_y}$, $\mathbf{b}_1 \in \mathbb{R}^{1 \times D_h}$, $\mathbf{b}_2 \in \mathbb{R}^{1 \times D_y}$.

- (c) Implement the forward and backward passes for a neural network with one sigmoid hidden layer. Fill in your implementation in `q1c_neural.py`. Sanity check your implementation with `python q1c_neural.py`.
- (d) GloVe (Global Vectors) embeddings are a type of word embeddings that represent words as vectors in a high-dimensional space, based on the co-occurrence statistics of words in a corpus. They are related to the skip-gram embeddings you saw in class in that they both aim to capture the semantic and syntactic relationships between words, but GloVe embeddings incorporate global corpus-level information in addition to local context information. In this section you will be using GloVe embeddings to represent the vocabulary. Use the neural network to implement a bigram language model in `q1d_neural_lm.py`. Use GloVe embeddings to represent the vocabulary (`data/lm/vocab.embeddings.glove.txt`). Implement the `lm_wrapper` function, that is used by `sgd` to sample the gradient, and the `eval_neural_lm` function that is used for model evaluation. Report the dev perplexity in your written solution. Don't forget to include `saved_params_40000.npy` in your submission zip!

2 Theoretical Inquiry of a Simple RNN Language Model

In this section we will perform a short theoretical analysis of a simple RNN language model, adapted from a paper by Tomas Mikolov, et al.¹. Formally, for every timestep t , the model is defined as follows:

$$\begin{aligned} \mathbf{e}^{(t)} &= \mathbf{x}^{(t)} \mathbf{L} \\ \mathbf{h}^{(t)} &= \text{sigmoid} \left(\mathbf{h}^{(t-1)} \mathbf{H} + \mathbf{e}^{(t)} \mathbf{I} + \mathbf{b}_1 \right) \\ \hat{\mathbf{y}}^{(t)} &= \text{softmax} \left(\mathbf{h}^{(t)} \mathbf{U} + \mathbf{b}_2 \right) \end{aligned} \quad (1)$$

where $\mathbf{h}^{(0)} \in \mathbb{R}^{D_h}$ is some initialization vector for the hidden layer and $\mathbf{x}^{(t)} \mathbf{L}$ is the product of \mathbf{L} with the one-hot vector $\mathbf{x}^{(t)}$ representing index of the current word. The parameters are:

$$\mathbf{L} \in \mathbb{R}^{|V| \times d} \quad \mathbf{H} \in \mathbb{R}^{D_h \times D_h} \quad \mathbf{I} \in \mathbb{R}^{d \times D_h} \quad \mathbf{b}_1 \in \mathbb{R}^{D_h} \quad \mathbf{U} \in \mathbb{R}^{D_h \times |V|} \quad \mathbf{b}_2 \in \mathbb{R}^{|V|} \quad (2)$$

where \mathbf{L} is the embedding matrix, \mathbf{I} is the input word weight matrix, \mathbf{H} is the hidden state weight matrix, \mathbf{U} is the output word transformation matrix, and \mathbf{b}_1 and \mathbf{b}_2 are biases. As for the dimensions, $|V|$ is the vocabulary size, d is the embedding dimension, and D_h is the hidden state dimension.

The output vector $\hat{\mathbf{y}}^{(t)} \in \mathbb{R}^{|V|}$ is a probability distribution over the vocabulary, and we optimize the cross-entropy loss:

$$J^{(t)}(\theta) = \text{CE}(\mathbf{y}^{(t)}, \hat{\mathbf{y}}^{(t)}) = - \sum_{i=1}^{|V|} y_i^{(t)} \log(\hat{y}_i^{(t)})$$

where $\mathbf{y}^{(t)}$ is the one-hot vector corresponding to the target word (which in our case is equal to $\mathbf{x}^{(t+1)}$). Note that $J^{(t)}(\theta)$ is a loss for a single timestep.

Compute the gradients for all model parameters at a single point in time (time-step) t :

$$\frac{\partial J^{(t)}}{\partial \mathbf{U}} \quad \frac{\partial J^{(t)}}{\partial \mathbf{L}_{\mathbf{x}^{(t)}}} \quad \frac{\partial J^{(t)}}{\partial \mathbf{I}} \Big|_{(t)} \quad \frac{\partial J^{(t)}}{\partial \mathbf{H}} \Big|_{(t)}$$

Where $\mathbf{L}_{\mathbf{x}^{(t)}}$ is the row of \mathbf{L} corresponding to the current input word $\mathbf{x}^{(t)}$ and $\Big|_{(t)}$ denotes the gradient for the appearance of that parameter at time t . (Equivalently, $\mathbf{h}^{(t-1)}$ is taken to be fixed, and you don't need to back-propagate to earlier time-steps just yet - you'll do that in part ??).

Additionally, compute the derivative with respect to the *previous* hidden layer value:²

$$\frac{\partial J^{(t)}}{\partial \mathbf{h}^{(t-1)}}$$

¹http://www.fit.vutbr.cz/research/groups/speech/publi/2010/mikolov_interspeech2010_IS100722.pdf.

You might Recognize Mikolov from <https://arxiv.org/abs/1301.3781>.

²For those of you who took Intro to ML, this derivative is also known as an “error term”, $\delta^{(t-1)}$.

3 Generating Shakespeare Using a Character-level Language Model

In this section we will train a language model and use it to generate text.

Follow the instructions, complete the code, and answer the questions from this Google Colab notebook³:
<https://colab.research.google.com/drive/1WIUACyCAgrPiuKzNBwXNChOzWrecLnCF?usp=sharing>

4 Perplexity

- (a) Show that perplexity calculated using the natural logarithm $\ln(x)$ is equal to perplexity calculated using $\log_2(x)$. i.e:

$$2^{-\frac{1}{M} \sum_{i=1}^M \log_2 p(s_i | s_1, \dots, s_{i-1})} = e^{-\frac{1}{M} \sum_{i=1}^M \ln p(s_i | s_1, \dots, s_{i-1})}$$

- (b) In this section you will be computing the perplexity of your previous trained models on two different passages. Please provide your results in the PDF file, as well as attach the code to your code files. The two different passages appear in the .zip file you've got. Their names are: `shakespeare_for_perplexity.txt` which contains a subset from the Shakespeare dataset, and `wikipedia_for_perplexity.txt` which contains a certain passage from Wikipedia. Please compute the perplexity of the bi-gram LM, and the model from section 3, on both these passages.
- (c) Try to explain the results you've got. Particularly, why there might be large gaps in perplexity, while looking on different passages.

5 Right-to-left vs left-to-right Estimation

Let x_0, x_1, \dots, x_n be any sentence, where x_0 is the start symbol and x_n is the end symbol. Prove that estimating $P(x_0, x_1, \dots, x_n)$ with a left-to-right count-based bi-gram model (which uses $P(x_i | x_{i-1})$) is equal to estimating it with a right-to-left count-based bi-gram model ($P(x_i | x_{i+1})$):

$$P(x_0 x_1 \dots x_n) = P(x_n) P(x_{n-1} | x_n) \dots P(x_0 | x_1) = P(x_0) P(x_1 | x_0) \dots P(x_n | x_{n-1}).$$

³Feel free to comment inside the notebook.